

Multitask Pointer Network for Discontinuous Constituent Parsing
An application to the German language

James Yu
jby21@cam.ac.uk

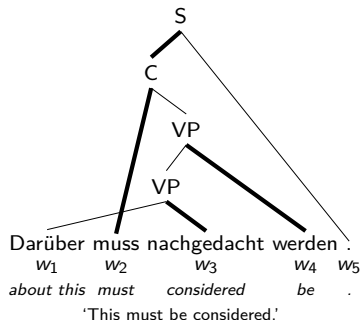
Faculty of Economics
University of Cambridge

15th January 2024

Why use a neural network to model German grammar?

- Constituent trees are a syntactic formalism representing phrasal hierarchy in a sentence.
- Free-order languages like German contain many grammatical discontinuities.
- Discontinuous representations introduce computational complexity but can be more valuable for downstream applications.
- Grammar-less neural network-based models have continually pushed the state of the art.
- My model is based on Fernández-González and Gómez-Rodríguez (2022), who propose an architecture based on pointer neural networks in a multi-task setting.
- I achieve state-of-the-art performance across several metrics.

Linguistic preliminaries: constituent trees



Let $w = (w_1, \dots, w_L)$ be a sentence.

Definition

- A *constituent tree* is a rooted tree whose leaves are the words $(w_i)_{i=1}^L$ and internal nodes are constituents satisfying some constraints.
- A *constituent* is a triple (Z, \mathcal{Y}, h) containing, respectively, its label, yield, and lexical head.
- A constituent is *discontinuous* if its yield is not contiguous.

Reduction to dependency parsing

Definition

A *dependency tree* is a rooted tree spanning the words in the sentence $(w_i)_{i=1}^L$. Each edge is labelled and connects a *head word* (parent) to a *dependency* (child).

Fernández-González and Martins (2015) show that constituent trees are isomorphic to dependency trees in which the edges contain information about constituent labels and attachment order.

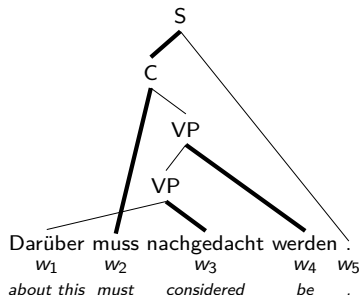


Figure: Constituent tree

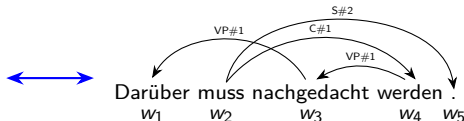
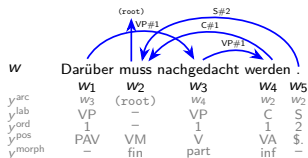


Figure: Dependency tree

Mathematical formalisation

- Regressor: sentence $(w_i)_{i=1}^L$.
- Regressand: $y_i = (y_i^{\text{arc}}, y_i^{\text{lab}}, y_i^{\text{ord}}, y_i^{\text{pos}}, y_i^{\text{morph}})$ for each $i = 1, \dots, L$.
- Bottom-up approach: think of *arcs* going from every child w_i to its parent $y_i^{\text{arc}} \in w \setminus w_i$.



Denote $y = (y_i)_{i=1}^L$ and with mild abuse of notation let $y_{<i} = (y_1, \dots, y_{i-1})$ for each $i = 2, \dots, L$ and $y_{<1} = 0$.

Assumption (conditional independence)

For each $i = 1, \dots, L$, the random variables $y_i^{\text{lab}}, y_i^{\text{ord}}, y_i^{\text{pos}}$ and y_i^{morph} are mutually independent conditional on $y_i^{\text{arc}}, y_{<i}$ and w .

We can decompose the conditional probability of y given w :¹

$$\begin{aligned}
 p_w(y) &= \prod_{i=1}^L p_w(y_i \mid y_{<i}) \\
 &= \prod_{i=1}^L \left\{ p_w(y_i^{\text{arc}} \mid y_{<i}) p_w(y_i^{\text{lab}} \mid y_i^{\text{arc}}, y_{<i}) \right. \\
 &\quad \cdot p_w(y_i^{\text{ord}} \mid y_i^{\text{arc}}, y_{<i}) p_w(y_i^{\text{pos}} \mid y_i^{\text{arc}}, y_{<i}) p_w(y_i^{\text{morph}} \mid y_i^{\text{arc}}, y_{<i}) \left. \right\}.
 \end{aligned}$$

¹For notational simplicity we have written $p_w(\cdot) \equiv p(\cdot \mid w)$ and $p_w(\cdot \mid \cdot) \equiv p(\cdot \mid \cdot, w)$.

Sequence-to-sequence setup

Given an input sentence $w = (w_i)_{i=1}^L$ we generate *embeddings* $\omega = (\omega_i)_{i=1}^L$, where

$$\omega_i = \mathbf{WordEmbed}(w_i) \oplus \mathbf{CharEmbed}(w_i) \oplus \mathbf{BertEmbed}(w_i).$$

- **WordEmbed** is a simple lookup table.
- **CharEmbed** is implemented using a CNN á la Chiu and Nichols (2016).
- BERT model pre-trained on German text by Chan et al. (2020).

Encoder: feed embeddings through a multi-layer bi-directional LSTM with skip-connections and dropout:

$$\mathbf{e} = (\mathbf{e}_i)_{i=0,\dots,L} = \mathbf{BiLSTM}(\omega).$$

(\mathbf{e}_0 is the initial state and represents the root pseudo-node.)

Decoder: feed embeddings through a single-layer uni-directional LSTM with dropout:

$$\mathbf{d} = (\mathbf{d}_i)_{i=1,\dots,L} = \mathbf{LSTM}((\mathbf{e}_i)_{i=1,\dots,L}).$$

(the initial state of the decoder is the final state of the encoder.)

Classification tasks: quadratic classifier

Building on Dozat and Manning (2016).

- Model conditional probabilities of y_i^{lab} , y_i^{ord} , y_i^{pos} and y_i^{morph} .

- Encoder and decoder are shared across tasks.

Example: part-of-speech classification.

$$e_i^{\text{pos}} = \text{MLP}_{\text{enc}}^{\text{pos}}(e_i); \quad d_j^{\text{pos}} = \text{MLP}_{\text{dec}}^{\text{pos}}(d_j).$$

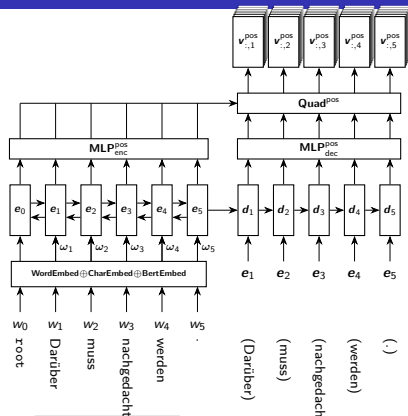
- Obtain class logits $v_{i,j}^{\text{pos}}$:

$$v_{i,j}^{\text{pos}} = \text{Quad}^{\text{pos}}(e_i^{\text{pos}}, d_j^{\text{pos}})$$

$$:= e_i^{\text{posT}} \mathbf{U}_{\text{h-d}}^{\text{pos}} d_j^{\text{pos}} + \boxed{e_i^{\text{posT}} \mathbf{U}_{\text{h-h}}^{\text{pos}} e_i^{\text{pos}}} + U_{\text{h}}^{\text{pos}} e_i^{\text{pos}} + \boxed{d_j^{\text{posT}} \mathbf{U}_{\text{d-d}}^{\text{pos}} d_j^{\text{pos}}} + U_{\text{d}}^{\text{pos}} d_j^{\text{pos}} + u^{\text{pos}}_{\text{bias}}.$$

Fixing child w_j , the vector $\text{softmax}(\mathbf{v}_{i,j}^{\text{pos}})$ can be interpreted as a probability distribution over its parts of speech conditional on having an arc to w_i :

$$p^{\text{pos}}(c \mid w_i, y_{<j}, w) = \text{softmax}(\mathbf{v}_{i,j}^{\text{pos}})_c.$$



The pointer network: quadratic attention mechanism

Building on Dozat and Manning (2016)'s bi-affine mechanism and drawing from Vinyals et al. (2015).

- Obtain dimension-reduced representations:

$$\mathbf{e}_i^{\text{arc}} = \text{MLP}_{\text{enc}_i}^{\text{arc}}(\mathbf{e}_i); \quad \mathbf{d}_j^{\text{arc}} = \text{MLP}_{\text{dec}_j}^{\text{arc}}(\mathbf{d}_j).$$

- Obtain *latent features* $\mathbf{v}_{i,j}^{\text{arc}}$:

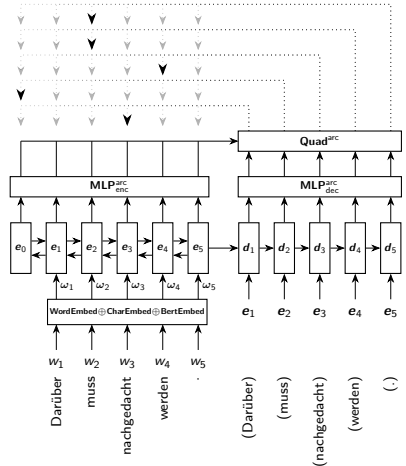
$$\begin{aligned} \mathbf{v}_{i,j}^{\text{arc}} &= \text{Quad}^{\text{arc}}(\mathbf{e}_i^{\text{arc}}, \mathbf{d}_j^{\text{arc}}) \\ &:= \mathbf{e}_i^{\text{arcT}} \mathbf{U}_{\text{h-d}}^{\text{arc}} \mathbf{d}_j^{\text{arc}} \\ &\quad + \mathbf{e}_i^{\text{arcT}} \mathbf{U}_{\text{h-h}}^{\text{arc}} \mathbf{e}_i^{\text{arc}} + U_{\text{h}}^{\text{arc}} \mathbf{e}_i^{\text{arc}} \\ &\quad + \mathbf{d}_j^{\text{arcT}} \mathbf{U}_{\text{d-d}}^{\text{arc}} \mathbf{d}_j^{\text{arc}} + U_{\text{d}}^{\text{arc}} \mathbf{d}_j^{\text{arc}} + \mathbf{u}_{\text{bias}}^{\text{arc}}. \end{aligned}$$

- Obtain *attention logits* $s_{i,j}^{\text{arc}}$:

$$s_{i,j}^{\text{arc}} = \mathbf{u}_{\text{agg}}^{\text{arcT}} \tanh(\mathbf{v}_{i,j}^{\text{arc}}).$$

- Fixing child w_j , the vector $\text{softmax}(\mathbf{s}_{:,j}^{\text{arc}})$ can be interpreted as a probability distribution over potential parents:

$$p^{\text{arc}}(w_i \mid y_{<j}, w) = \text{softmax}(\mathbf{s}_{:,j}^{\text{arc}})_i.$$



Training

We find \hat{p} using *SGD* with *Nesterov momentum*, finding a suitably low *cross-entropy* between the model p and the empirical distribution present in the dataset $(w^n, y^n)_{n=1}^N$:

$$\begin{aligned} & - \sum_{n=1}^N \log \prod_{i=1}^{L_n} \left\{ p_w(y_i^{n\text{arc}} | y_{<i}^n) p_w(y_i^{n\text{lab}} | y_i^{n\text{arc}}, y_{<i}^n) p_w(y_i^{n\text{ord}} | y_i^{n\text{arc}}, y_{<i}^n) \right. \\ & \quad \cdot p_w(y_i^{n\text{pos}} | y_i^{n\text{arc}}, y_{<i}^n) p_w(y_i^{n\text{morph}} | y_i^{n\text{arc}}, y_{<i}^n) \end{aligned}$$
$$= \text{loss}^{\text{arc}} + \text{loss}^{\text{lab}} + \text{loss}^{\text{ord}} + \text{loss}^{\text{pos}} + \text{loss}^{\text{morph}}.$$

Inference

Given $w = (w_i)_{i=1}^L$, estimate $y = (y_i)_{i=1}^L$ by maximising the estimated conditional probability:

$$\hat{y} = \arg \max_y \left\{ \prod_{i=1}^L \left\{ \hat{p}_w(y_i^{\text{arc}} | y_{<i}) \hat{p}_w(y_i^{\text{lab}} | y_i^{\text{arc}}, y_{<i}) \right. \right. \\ \cdot \hat{p}_w(y_i^{\text{ord}} | y_i^{\text{arc}}, y_{<i}) \hat{p}_w(y_i^{\text{pos}} | y_i^{\text{arc}}, y_{<i}) \hat{p}_w(y_i^{\text{morph}} | y_i^{\text{arc}}, y_{<i}) \left. \right\} \Bigg\}.$$

The feasible region is very large, so the maximisation is approximated via *beam search*.

Model evaluation

The TIGER treebank is a widely-used corpus of $\sim 50\,000$ constituent trees. Its main textual basis is the *Frankfurter Rundschau*.

- 97 % of the dataset is usable as training examples, with train/dev/test split of 80/10/10 %.
- PARSEVAL metrics initially proposed by Black et al. (1991):

$$P = \frac{\# \text{ of correct constituents in prediction}}{\# \text{ of total constituents in prediction}}; \quad R = \frac{\# \text{ of correct constituents in prediction}}{\# \text{ of total constituents in reference}}.$$

Model	F1	Disc. F1
Coavoux et al. (2019)	82.7	55.9
Corro (2020)	90.0	62.1
F.-González & G.-Rodríguez (2022)	89.8	71.0
Chen & Komachi (2023)	89.6	70.9
This work	90.59	84.74

Table: Comparison of overall F1-score (%) and F1-score measured only on discontinuous constituents (disc. F1). Calculated using disco-dop (Van Cranenburgh et al., 2016) as standard practice. All models configured with BERT.

Model	pos	morph (avr)
Müller et al. (2013)	98.20	98.27
Schnabel & Schütze (2014)	97.50	97.76
Kondratyuk et al. (2018)	98.58	98.97
This work	99.21	99.60

Table: Comparison of part of speech (pos) and morphology accuracies (%). Morphology accuracies are the average of accuracies for case, degree, gender, mood, number, person and tense.

Appendix: Mathematical justification for quadratic classifier

Notation has been simplified for presentational clarity. Single c is multinoulli and

$$\begin{bmatrix} \mathbf{e}_i \\ \mathbf{d}_j \end{bmatrix} \mid c \sim \mathcal{N} \left(\begin{bmatrix} \boldsymbol{\mu}_c \\ \boldsymbol{\phi}_c \end{bmatrix}, \begin{bmatrix} A_c & Q_c^\top \\ Q_c & B_c \end{bmatrix}^{-1} \right) \implies \mathbf{d}_j \mid c \sim \mathcal{N}(\boldsymbol{\phi}_c, P_c^{-1}).$$

The conditional log-probability of c is the following affine quadratic form:


$$\begin{aligned} \log p(c \mid \mathbf{e}_i, \mathbf{d}_j) &= k_c - \frac{1}{2} \left((\mathbf{e}_i - \boldsymbol{\mu}_c) + A_c^{-1} Q_c^\top (\mathbf{d}_j - \boldsymbol{\phi}_c) \right)^\top A_c \left((\mathbf{e}_i - \boldsymbol{\mu}_c) + A_c^{-1} Q_c^\top (\mathbf{d}_j - \boldsymbol{\phi}_c) \right) \\ &\quad - \frac{1}{2} (\mathbf{d}_j - \boldsymbol{\phi}_j)^\top P_c (\mathbf{d}_j - \boldsymbol{\phi}_j) - \log p(\mathbf{e}_i, \mathbf{d}_j) \\ &= - \boxed{\mathbf{e}_i^\top Q_c^\top \mathbf{d}_j} - \boxed{\frac{1}{2} \mathbf{e}_i^\top A_c \mathbf{e}_i} - \boxed{\frac{1}{2} \mathbf{d}_j^\top (P_c + Q_c A_c^{-1} Q_c^\top) \mathbf{d}_j} + \boxed{(\boldsymbol{\mu}_c^\top A_c + \boldsymbol{\phi}_c^\top Q_c) \mathbf{e}_i} \\ &\quad + \boxed{(\boldsymbol{\phi}_c^\top Q_c A_c^{-1} Q_c^\top \boldsymbol{\phi}_c^\top P_c + \boldsymbol{\mu}_c^\top Q_c^\top) \mathbf{d}_j} \\ &\quad + \boxed{-\boldsymbol{\mu}_c^\top Q_c^\top \boldsymbol{\phi}_c - \frac{1}{2} \boldsymbol{\mu}_c^\top A_c \boldsymbol{\mu}_c - \frac{1}{2} \boldsymbol{\phi}_c^\top (P_c + Q_c A_c^{-1} Q_c^\top) \boldsymbol{\phi}_c + k_c} \\ &= \text{cth row of } \mathbf{e}_i^\top \mathbf{U}_{h-d} \mathbf{d}_j + \mathbf{e}_i^\top \mathbf{U}_{h-h} \mathbf{e}_i + \mathbf{d}_j^\top \mathbf{U}_{d-d} \mathbf{d}_j + \mathbf{U}_h \mathbf{e}_i + \mathbf{U}_d \mathbf{d}_j + \mathbf{u}_{\text{bias}}. \end{aligned}$$

$-\log p(\mathbf{e}_i, \mathbf{d}_j)$
 constant in c
 so throw away

Memory-efficient implementation: assuming conditional normality:

$$v_{i,j}^c = k^c - \|W_1^c \mathbf{e}_i + W_2^c \mathbf{d}_j + \mathbf{w}_3^c\|_2^2.$$

Bibliography I

- 
- Black, E., Abney, S., Flickinger, D., Gdaniec, C., Grishman, R., Harrison, P., Hindle, D., Ingria, R., Jelinek, F., Klavans, J. L., et al. (1991). A procedure for quantitatively comparing the syntactic coverage of english grammars. *Speech and Natural Language: Proceedings of a Workshop Held at Pacific Grove, California, February 19-22, 1991*.
- 
- Chan, B., Möller, T., Pietsch, M., & Soni, T. (2020). *German BERT (bert-base-german-cased)*. Retrieved September 1, 2023, from <https://huggingface.co/bert-base-german-cased>
- 
- Chen, Z., & Komachi, M. (2023). Discontinuous combinatory constituency parsing. *Transactions of the Association for Computational Linguistics*, 11, 267–283.
- 
- Chiu, J. P., & Nichols, E. (2016). Named entity recognition with bidirectional lstm-cnns. *Transactions of the association for computational linguistics*, 4, 357–370.
- 
- Coavoux, M., Crabbé, B., & Cohen, S. B. (2019). Unlexicalized transition-based discontinuous constituency parsing. *Transactions of the Association for Computational Linguistics*, 7, 73–89.
- 
- Corro, C. (2020). Span-based discontinuous constituency parsing: A family of exact chart-based algorithms with time complexities from $O(n^6)$ down to $O(n^3)$. *Empirical Methods in Natural Language Processing*.
- 
- Dozat, T., & Manning, C. D. (2016). Deep biaffine attention for neural dependency parsing. *International Conference on Learning Representations*.

Bibliography II



Eger, S., Gleim, R., & Mehler, A. (2016). Lemmatization and morphological tagging in german and latin: A comparison and a survey of the state-of-the-art. *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*, 1507–1513.



Fernández-González, D., & Gómez-Rodríguez, C. (2022). Multitask pointer network for multi-representational parsing. *Knowledge-Based Systems*, 236, 107760.



Fernández-González, D., & Martins, A. F. (2015). Parsing as reduction. *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, 1523–1533.



Gleim, R., Eger, S., Mehler, A., Uslu, T., Hemati, W., Lücking, A., Henlein, A., Kahlsdorf, S., & Hoenen, A. (2019). A practitioner's view: A survey and comparison of lemmatization and morphological tagging in german and latin. *Journal of Language Modelling*, 7(1), 1–52.



Kondratyuk, D., Gavenčiak, T., Straka, M., & Hajič, J. (2018). Lemmatag: Jointly tagging and lemmatizing for morphologically-rich languages with brnns. *arXiv preprint arXiv:1808.03703*.



Müller, T., Schmid, H., & Schütze, H. (2013). Efficient higher-order crfs for morphological tagging. *Proceedings of the 2013 conference on empirical methods in natural language processing*, 322–332.



Schnabel, T., & Schütze, H. (2014). FLORS: Fast and simple domain adaptation for part-of-speech tagging. *Transactions of the Association for Computational Linguistics*, 2, 15–26.



Van Cranenburgh, A., Scha, R., & Bod, R. (2016). Data-oriented parsing with discontinuous constituents and function tags. *Journal of Language Modelling*, 4(1), 57–111.



Vinyals, O., Fortunato, M., & Jaitly, N. (2015). Pointer networks. *Advances in neural information processing systems*, 28.